

Professional Summary:

- Around **8** years of professional Information Technology experience in **Hadoop** and **Linux Administration** activities such as installation, configuration and maintenance of systems/clusters.
- Having extensive experience in **Linux Administration & Big Data Technologies** as a **Hadoop Administration**.
- Hands on experience in **Hadoop** Clusters using **Horton works** (HDP), **Cloudera** (CDH3, CDH4), oracle big data and **Yarn** distributions platforms.
- Possessing skills in Apache **Hadoop**, **Map-Reduce**, **Pig**, Impala, **Hive**, **base**, **Zookeeper**, **Sqoop**, **Flume**, **OOZIE**, **Kafka**, **storm**, **Spark**, **Java Script**, and **J2EE**.
- Experience in deploying and managing the multi-node development and production **Hadoop** cluster with different **Hadoop** components (**Hive**, **Pig**, **Sqoop**, **Oozie**, **Flume**, **HCatalog**, **HBase**, **Zookeeper**) using **Horton works Ambari**.
- Good experience in creating various database objects like tables, stored procedures, functions, and triggers using **SQL**, **PL/SQL** and **DB2**.
- Used **Apache Falcon** to support Data Retention policies for **HIVE/HDFS**.
- Experience in Configuring **Name-node High availability** and Name-node Federation and depth knowledge on **Zookeeper** for cluster coordination services.
- Experience on Design, configure and manage the backup and disaster recovery for **Hadoop** data.
- Experience in administering **Tableau** and **Green Plum** databases instances in various environments.
- Experience in administration of **Kafka** and **Flume** streaming using **Cloudera Distribution**.
- Hands on experience in analyzing Log files for **Hadoop** and eco system services and finding root cause.
- Experience on Commissioning, Decommissioning, Balancing, and Managing Nodes and tuning server for optimal performance of the cluster.
- Experience in importing and exporting the data using **Sqoop** from **HDFS** to Relational Database systems/mainframe and vice-versa.
- Worked on **NoSQL** databases including **HBase**, **Cassandra** and **MongoDB**.
- Designing and implementing security for **Hadoop** cluster with Kerberos secure authentication.
- Hands on experience on **Nagios** and **Ganglia** tool for cluster monitoring system.
- Experience in scheduling all **Hadoop/Hive/Sqoop/HBase** jobs using **Oozie**.
- Experience in **HDFS** data storage and support for running **map-reduce jobs**.
- Experience in Installing Firmware Upgrades, kernel patches, systems configuration, performance tuning on **Unix/Linux systems**.
- Expert in Linux Performance monitoring, kernel tuning, Load balancing, health checks and maintaining compliance with specifications.
- Hands on experience in **Zookeeper** and **ZKFC** in managing and configuring in **Name Node** failure scenarios.
- Team Player with good communication and interpersonal skills and goal oriented approach to problem solving issues.

Technical Skills:

Big Data Technologies	Hadoop, HDFS, Map Reduce, Yarn, Hive, Pig, Sqoop, HBase, Flume, Oozie, Spark, Zookeeper.
Hadoop Platforms	Horton works and Cloudera, Apache Hadoop.
Networking Concepts	OSI Model, TCP/IP, UDP, IPV4, DHCP & DNS.
Programming Languages	Java, UNIX shell scripting and Bash.
Operating Systems	Linux (CentOS, Ubuntu, Red Hat), Windows, UNIX and Mac OS-X.
Database/ETL	Oracle, Cassandra, DB2, MS-SQL Server, MySQL, MS-Access, HBase, MongoDB, Teradata.

Narendra Reddy

Email: nbommaka.43@gmail.com

Phone: 512-466-3550

XML Languages	XML, DTD, XML Schema, XPath.
Monitoring and Alerting	Nagios, Ganglia, Cloudera Manager, Ambari.

Education: Bachelor of Technology (Information Technology) from SCSVMV University, TN, INDIA.

Professional Experience:

Role: Hadoop Administrator.
Client: T-Mobile Seattle, WA.

From: Aug'17 to Present.

Responsibilities:

- Working on Hadoop Horton works (HDP 2.5.3) distribution which managed services viz. HDFS, MapReduce2, Hive, Pig, HBase, Sqoop, Flume, Spark, Ambari Metrics, Zookeeper, Falcon and oozie etc. for 4 cluster ranges from LAB, DEV, QA to PROD.
- Monitor Hadoop cluster connectivity and security on Ambari monitoring system.
- Led the installation, configuration and deployment of product software's on new edge nodes that connect and contact Hadoop cluster for data acquisition.
- Responsible for Cluster maintenance, Monitoring, commissioning and decommissioning Data nodes, Troubleshooting, Manage and review data backups, Manage & review log files.
- Day to day responsibilities includes solving developer issues, deployments moving code from one environment to other environment, providing access to new users and providing instant solutions to reduce the impact and documenting the same and preventing future issues.
- Interacting with HDP support and log the issues in portal and fixing them as per the recommendations.
- Using Flume and Spool directory loading the data from local system to HDFS.
- Retrieved data from HDFS into relational databases with Sqoop.
- Extending the functionality of Hive and Pig with custom UDF's and UDAF's.
- Worked on analyzing Hadoop cluster and different big data analytic tools including Pig, HBase database and Sqoop.
- Commissioned and Decommissioned nodes on Hadoop cluster on Red Hat LINUX.
- Involved in loading data from LINUX file system to HDFS.
- Experience in configuring the Storm in loading the data from MYSQL to HBASE using jms.
- Integrated Kerberos into Hadoop to make cluster more strong and secure from unauthorized users.
- Exported the analyzed data to the relational databases using Sqoop for visualization and to generate reports for the BI team.
- Installed Oozie workflow engine to run multiple Hive and pig jobs.
- Analyzed large amounts of data sets to determine optimal way to aggregate and report on it.

Environment: Hadoop, HDFS, Map Reduce, HIVE, PIG, FLUME, OOZIE, Sqoop, Eclipse, Horton works, Ambari, Redhat, MYSQL.

Role: Hadoop Administrator
Client: AMD, Austin, TX.

From: Nov'15 to July'17.

Responsibilities:

- Installed/Configured/Maintained Apache Hadoop and Cloudera Hadoop clusters for application development and Hadoop tools like Hive, Pig, HBase, Zookeeper and Sqoop.
- Worked on Providing User support and application support through remedy ticket management system on Hadoop Infrastructure.
- Installed and configured Hadoop cluster in Development, Testing and Production environments.

Narendra Reddy

Email: nbommaka.43@gmail.com

Phone: 512-466-3550

- Performed both major and minor upgrades to the existing CDH cluster.
- Installation of various Hadoop Ecosystems and Hadoop Daemons.
- Installed and configured Flume agents with well-defined sources, channels and sinks.
- Configured safety valve to create active directory filters to sync the LDAP directory for Hue.
- Understanding the existing Enterprise data warehouse set up and provided design and architecture suggestion converting to Hadoop using Map Reduce, HIVE, SQOOP and Pig Latin.
- Implemented Name Node backup using NFS. This was done for High availability.
- Worked on importing and exporting data from Oracle and DB2 into HDFS and HIVE using Sqoop.
- Wrote shell scripts for rolling day-to-day processes and it automated using Crontab.
- Collected the logs data from web servers and integrated in to HDFS using Flume.
- Implemented FIFO schedulers on the Job tracker to share the resources of the Cluster for the Map Reduce jobs given by the users.
- Involved in Data model sessions to develop models for HIVE tables.

Environment: Apache Hadoop, CDH4, Hive, Hue, Pig, HBase, Map Reduce, Sqoop, Redhat, CentOS.

Role: Hadoop Administrator

From: Dec'14 to Oct'15.

Client: Katch LLC, Los Angeles, CA.

Responsibilities:

- Installed, configured and maintained Apache Hadoop clusters for application development and Hadoop tools like **Hive, Pig, HBase, Zookeeper and Sqoop**.
- Managing and scheduling Jobs on Hadoop Clusters using Apache, **Cloudera (CDH3, CDH4)** distributions.
- Worked on importing and exporting data from Oracle and DB2 into **HDFS using Sqoop**.
- Developed **PIG Latin scripts** to extract the data from the web server output files to load into **HDFS**.
- Implemented Name Node backup using **NFS**. This was done for High availability.
- Implemented Fair schedulers on the Job tracker to share the resources of the Cluster for the **Map Reduce** jobs given by the users.
- Configured custom interceptors in Flume agents for replicating and multiplexing data into multiple sinks.
- Worked on NoSQL databases including HBase and **MongoDB**.
- Setting up automated 24*7 monitoring and escalation infrastructure for Hadoop cluster using Nagios and Ganglia.
- Extensive experience in data analysis using tools like sync sort and **HZ** along with Shell Scripting and UNIX.
- Handle the data exchange between HDFS& Web Applications and databases using Flume and Sqoop.
- Experienced in developing Map Reduce programs using Apache Hadoop for working with Big Data.
- Expertise in working with different databases likes Oracle, MS-SQL Server, **PostgreSQL** and MS Access 2012 along with exposure to Hibernate for mapping an object-oriented domain model to a traditional relational database.
- Good understanding of Scrum methodologies, Test Driven Development and continuous integration.

Environment: Cloudera, CDH 4.4, and CDH 3, Cloudera manager, Sqoop, Flume, Hive, HQL, Pig, RHEL, Cent OS, Oracle, MS-SQL, Zookeeper, Oozie, Map Reduce, Apache Hadoop 1.x, PostgreSQL, Ganglia and Nagios.

Role: Linux Administrator

From: Sept'12 to Oct'13.

Client: Unitech Ltd, India.

Responsibilities:

- Installation, Configuration, Up gradation and administration of Windows, Sun Solaris, Red Hat Linux and Solaris.
- Linux and Solaris installation, administration and maintenance.
- User account management, managing passwords setting up quotas and support.
- Worked on Linux Kick-start OS integration, DDNS, DHCP, SMTP, Samba, NFS, FTP, SSH, and LDAP integration.
- Network traffic control, IPSec, Quos, VLAN, Proxy, Radius integration on Cisco Hardware via Red Hat Linux Software.
- Installation and configuration of MySQL on Windows Server nodes.
- Responsible for configuring and managing Squid server in Linux and Windows.
- Configuration and Administration of **NIS** environment.
- Package and Patch management on **Linux servers**.
- Worked on Logical volume manager to create file systems as per user and database requirements.
- Data migration at Host level using **Red Hat LVM**, **Solaris LVM**, and **Veritas Volume Manager**.
- Expertise in establishing and documenting the procedures to ensure data integrity including system fail-over and backup/recovery in **AIX** operating system.
- Managed **100 + UNIX servers** running **RHEL**, **HPUX** on **Oracle HP**.
- **Solaris Disk Mirroring (SVM)**, **ZONE** installation and configuration
- Escalating issues accordingly, managing team efficiently to achieve desired goals.

Environment: Linux, TCP/IP, LVM, RAID, Networking, Security, user management, MySQL.

Role: Linux/System Administrator

From: June'11 to Aug'12

Client: Scepter Technologies, India

Responsibilities:

- Installing / configuring RHEL (Redhat Enterprise Linux) .
- Supporting Web Logic servers and deployments.
- Supporting VMware installations.
- Preparing Operational Run books.
- Supporting Open Text CMS architecture.
- Redhat Package Management.
- Tomcat / JBOSS installation and support.
- Web Seal and DB2 - Basic knowledge.
- Hands on experience on HP Data protector backup / restore utility.
- Search engine support of Verity and Google analytics.
- Installing / Supporting Jira / confluence.
- Supporting Linux environment for disaster recovery environment.
- Installing and maintaining the Linux servers.
- Installed Cent OS using Pre-Execution environment boot and Kick start method on multiple servers.
- Monitoring System Metrics and logs for any problems.
- Running Crontab to back up data.
- Maintaining the MYSQL server and Authentication to required users for databases.
- Creating and managing Logical volumes.
- Installing and updating packages using YUM.

Environment: MYSQL, PHP, DB2, APACHE, MYSQL WORKBENCH, TOAD, LINUX.
